

Explanation of Gradient Descent Algorithm in Linear Regression Implementation

In my linear regression implementation file, the gradient descent algorithm in the ‘fit_linear_regression’ function may seem as if it was pulled out of thin air. For that reason, I have decided to include this document explaining the mathematics behind it, without going into too much detail and working.

The cost function of the univariate linear regression algorithm is:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$$

Where $h_\theta(x) = \theta_0 + \theta_1 x$.

The gradient descent algorithm is defined by:

Repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
(Simultaneously update j = 0 and j = 1)
}

After calculating the required partial derivatives and subbing into the algorithm, the gradient descent for univariate linear regression is given by:

Repeat until convergence {
 $\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})$
 $\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$
}

Which is vectorized in the gradient descent algorithm part of the my implementation.